

# Processes and Threads

The present slides are mainly adapted from «*Operating Systems: Internals and Design Principles*» 6/E by William Stallings (Chapter 4). Some materials are obtained from the POSIX threads Programming tutorial by Blaise Barney.


*Sistemi di Calcolo 2*

*Instructor: Riccardo Lazzeretti*

*Special thanks to: Daniele Cono D'Elia, Leonardo Aniello, Roberto Baldoni*



# Roadmap

- 
- Processes: `fork()`, `wait()`
  - Threads: resource ownership and execution
  - Case study:
    - Pthreads
  - Symmetric multiprocessing (SMP)





# What is a process?

- The **process** is a dynamic entity loaded on RAM memory generated by a program: more precisely, it is a sequence of activities (task) controlled by a program (scheduler) that takes place on a processor typically under the management or supervision of the respective operating system.





# Role of Processes

- Most requirements that an OS must meet can be expressed w.r.t. processes:
  - Interleaved execution
  - Resource allocation and policies
  - User creation of processes and inter-process communication





# Process Elements


- A process is comprised of:
  - Program code (possibly shared)
  - A set of data
  - A number of attributes describing the state of the process *during execution*






# Process Elements

- While the process is running it has a number of elements including
  - Identifier
  - State
  - Priority
  - Program counter
  - Memory pointers
  - Context data
  - I/O status information
  - Accounting information





# Process Control Block

- Contains the process elements
- Created and managed by the operating system
- Allows support for multiple processes




Figure 3.1 Simplified Process Control Block

# Unix system calls

## Creating new Processes

---

`fork( ) - wait( ) - exit( )`

*Credits: Mirela Damian, Allan Gottlieb*

*Class notes: <https://cs.nyu.edu/~gottlieb/courses/os202/class-notes.html>*




# How To Create New Processes?

---

## □ Underlying mechanism


- A process runs **fork** to create a child process
- Parent and children execute concurrently
- Child process is a duplicate of the parent process



# Process Creation

---

- After a **fork**, both parent and child keep running, and each can fork off other processes.
- A **process tree** results. The root of the tree is a special process created by the OS during startup.
- A process can *choose* to wait for children to terminate. For example, if C issued a **wait()** system call, it would block until G finished.



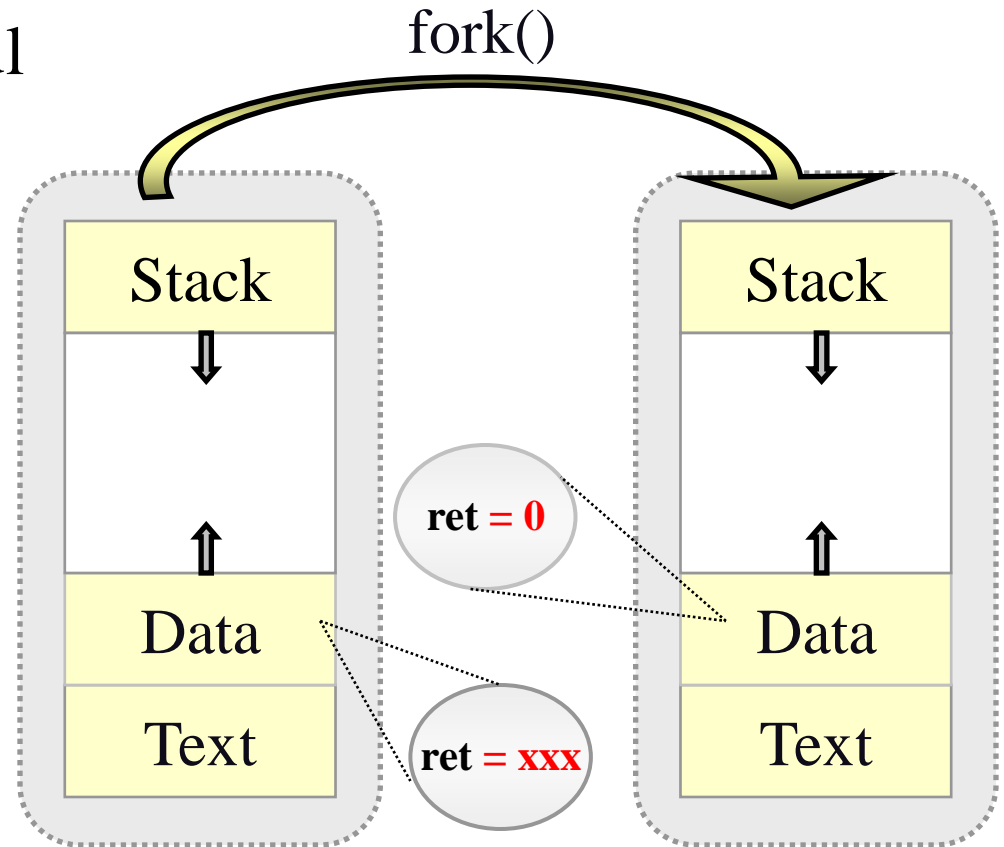
# Bootstrapping

---

- When a computer is switched on or reset, there must be an initial program that gets the system running
- This is the bootstrap program
  - Initialize CPU registers, device controllers, memory
  - Load the OS into memory
  - Start the OS running
- OS starts the first process (such as “init”)
- OS waits for some event to occur
  - Hardware interrupts or software interrupts (traps)

# Fork System Call

- Current process split into 2 processes: parent, child
- Returns -1 if unsuccessful
- Returns 0 in the child
- Returns the child's identifier in the parent

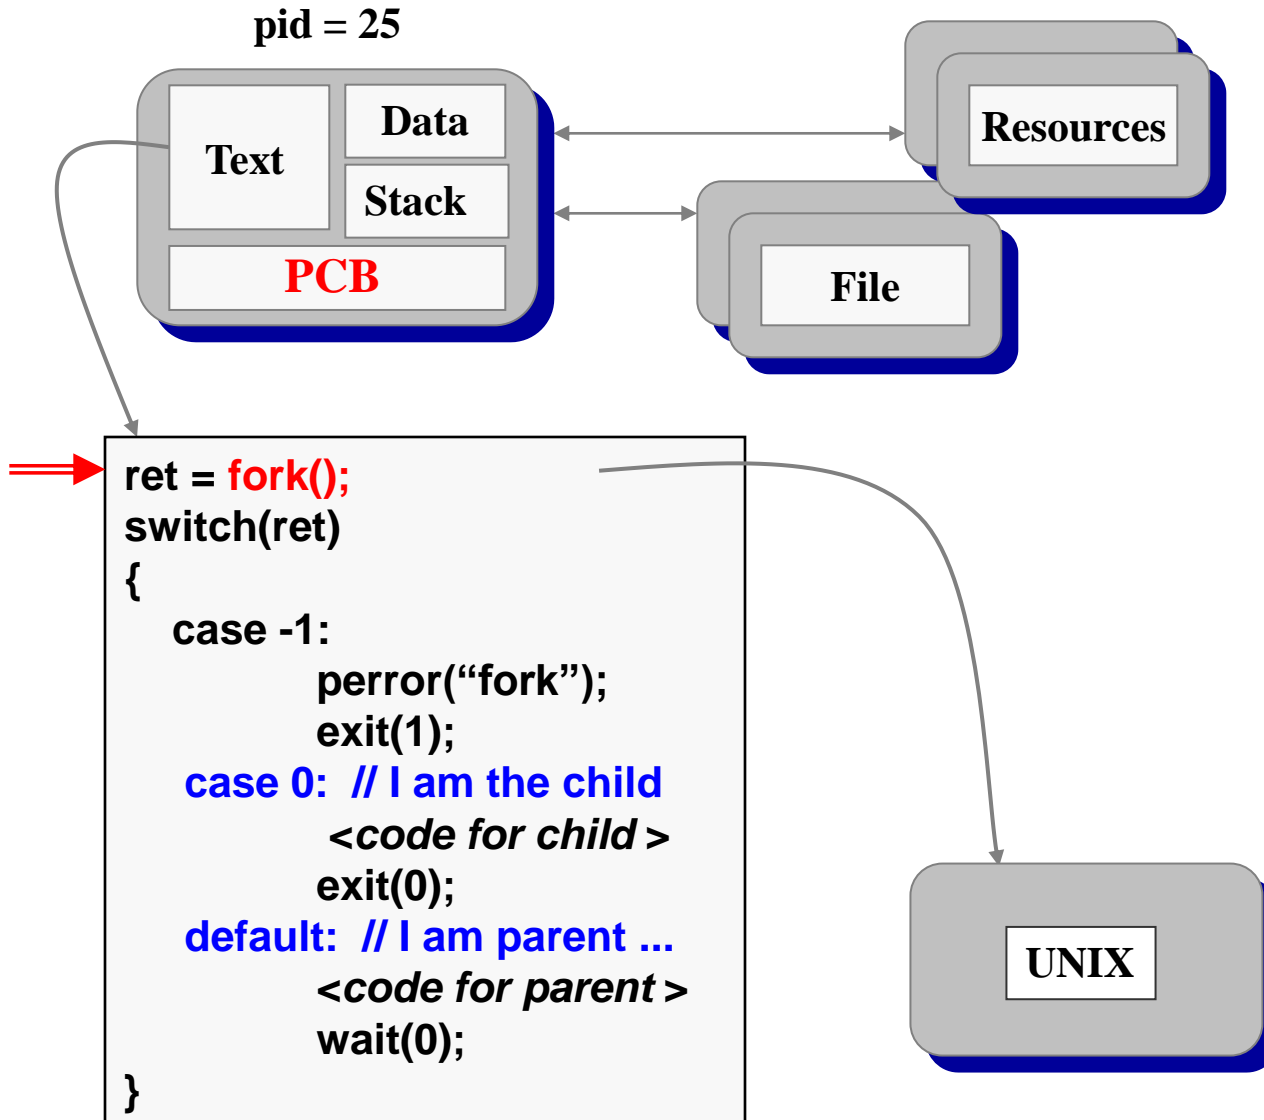


# Fork System Call

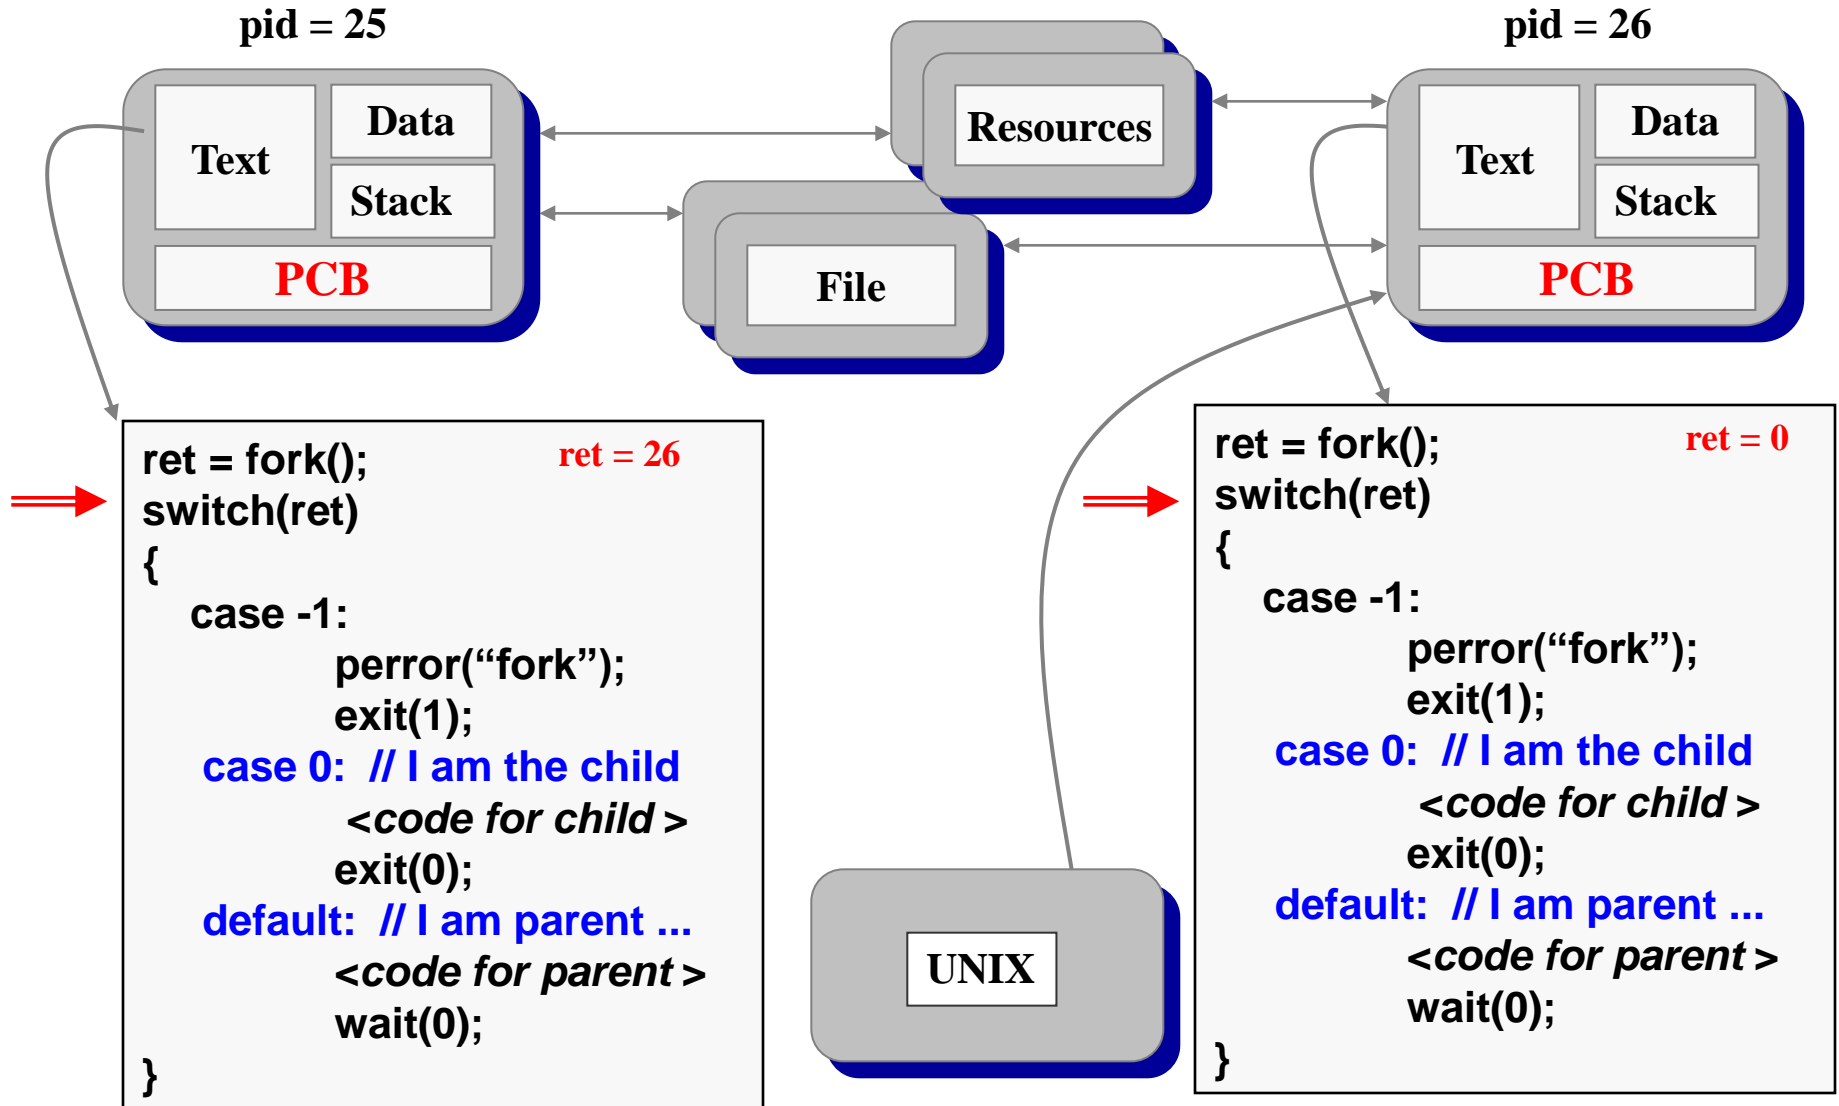
---

- The child process inherits from parent
  - identical copy of memory
  - CPU registers
  - all files that have been opened by the parent
- Execution proceeds **concurrently** with the instruction following the fork system call
- The execution context (PCB) for the child process is a copy of the parent's context at the time of the call

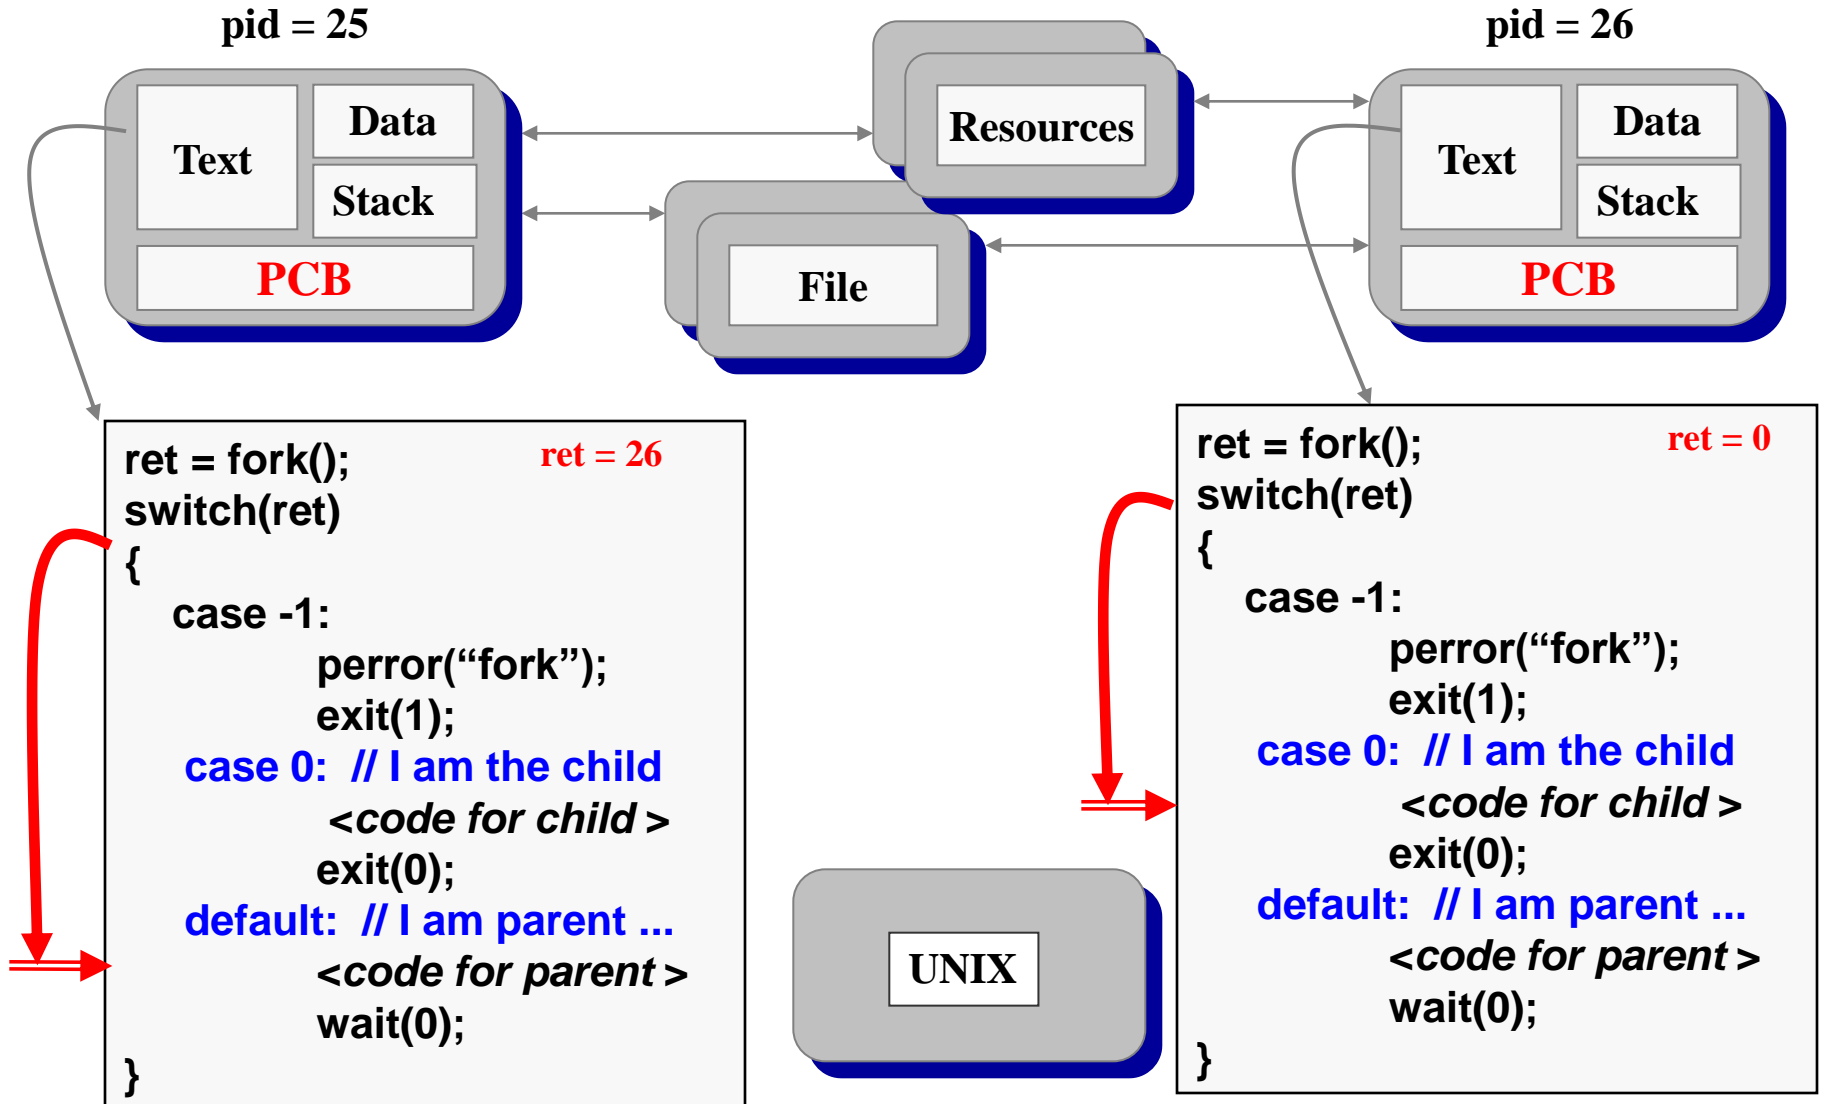
# How fork Works (1)



# How fork Works (2)

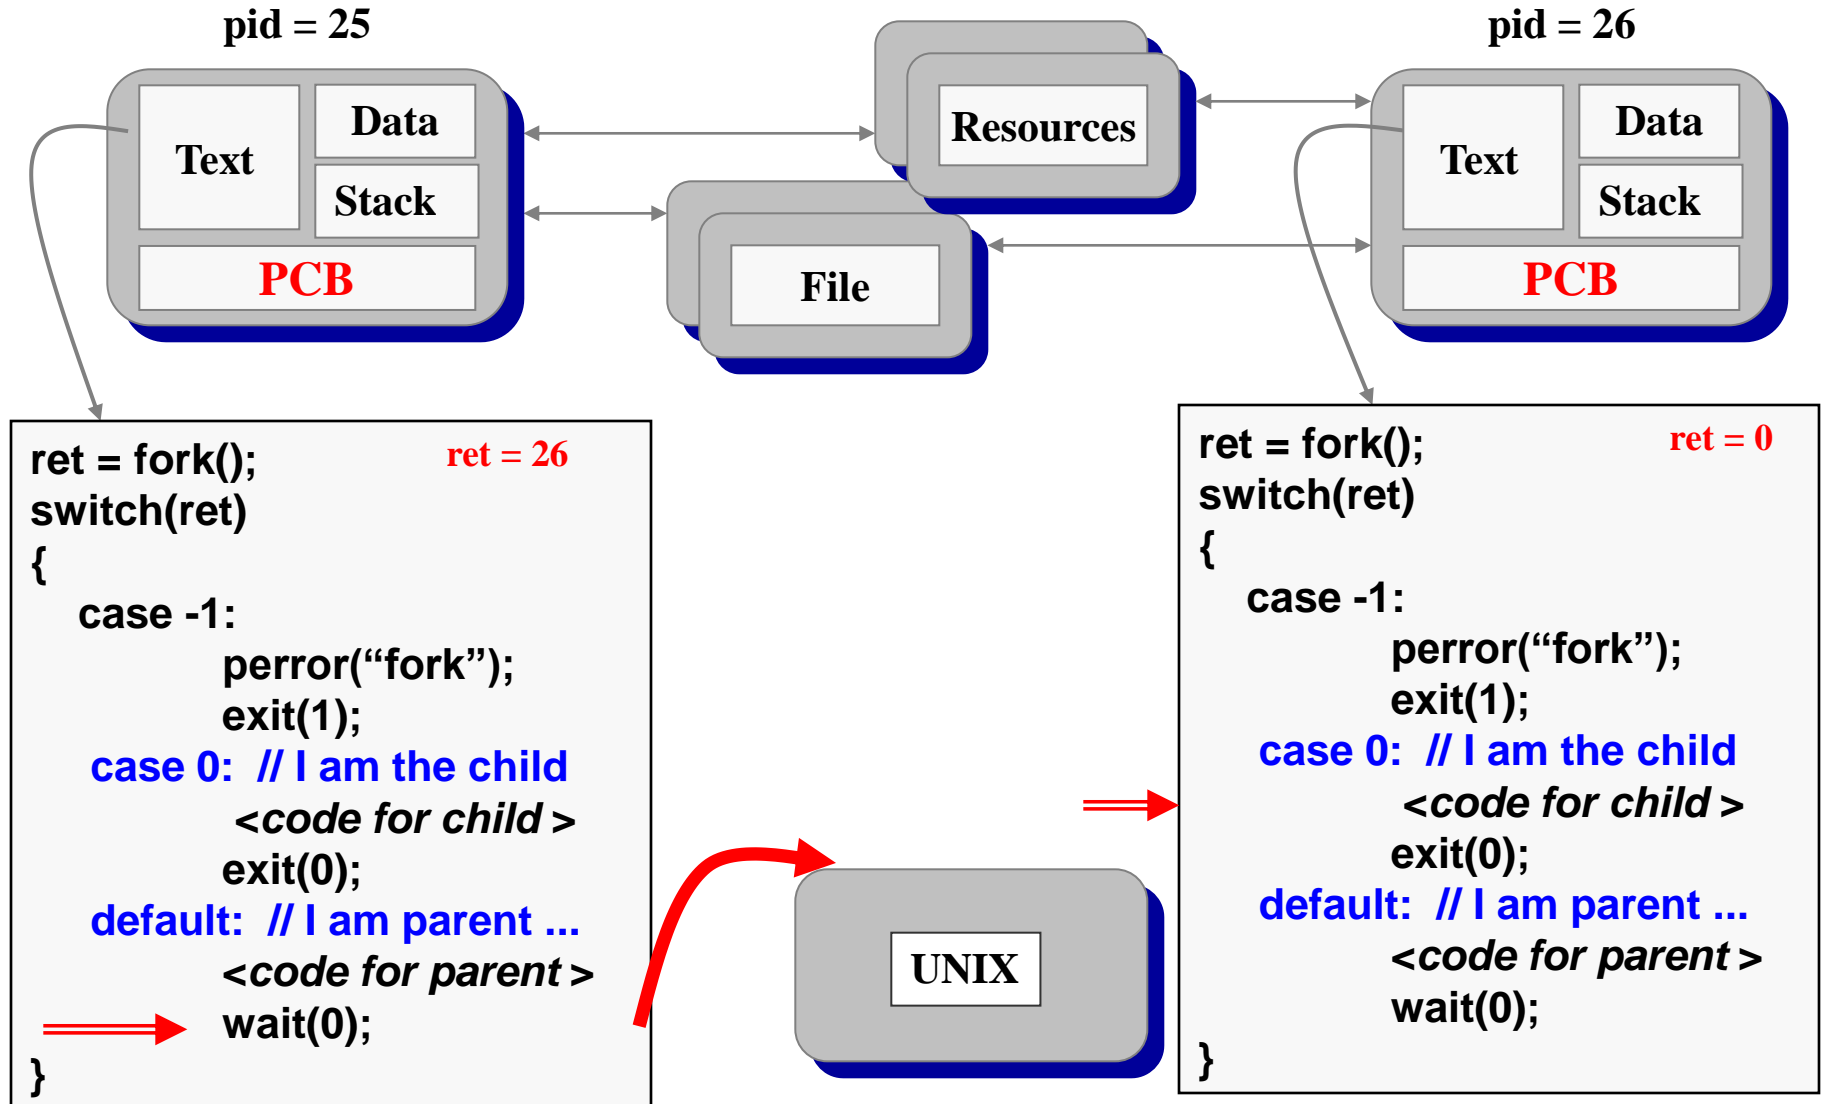


# How fork Works (3)

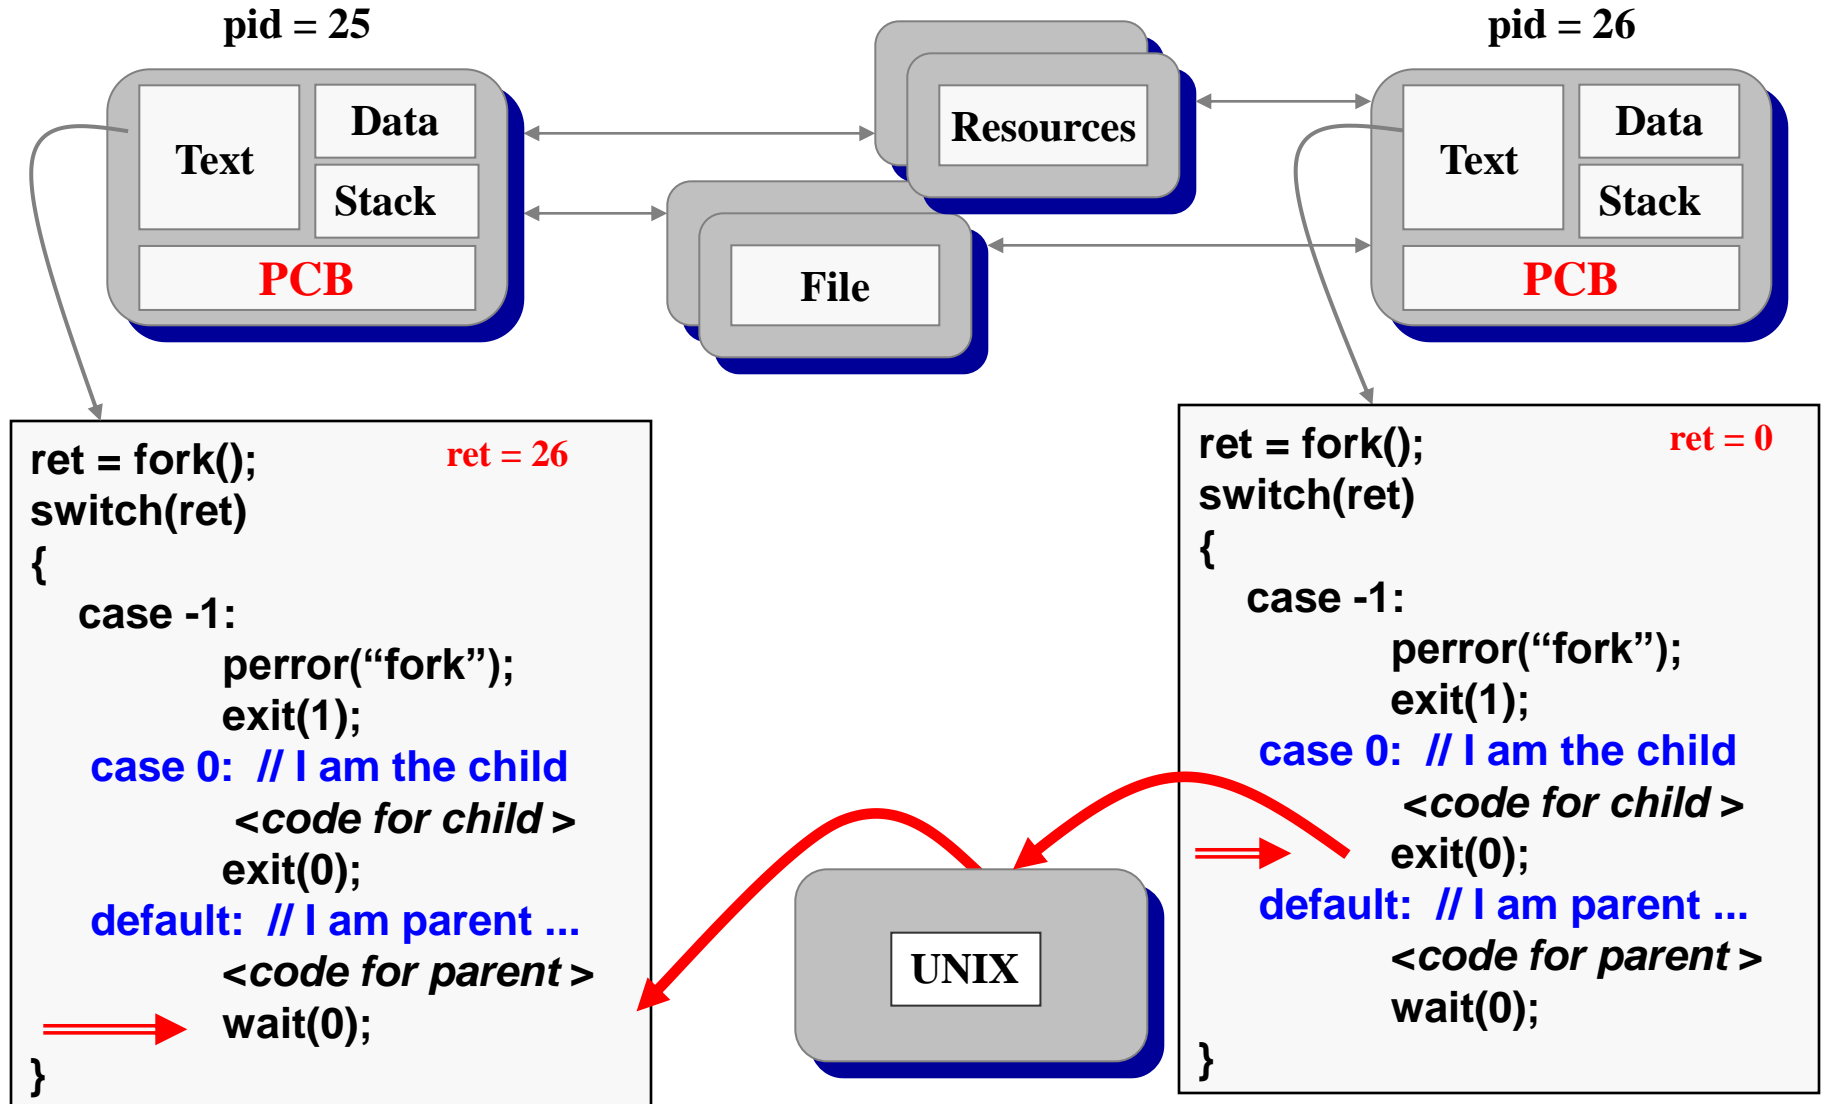




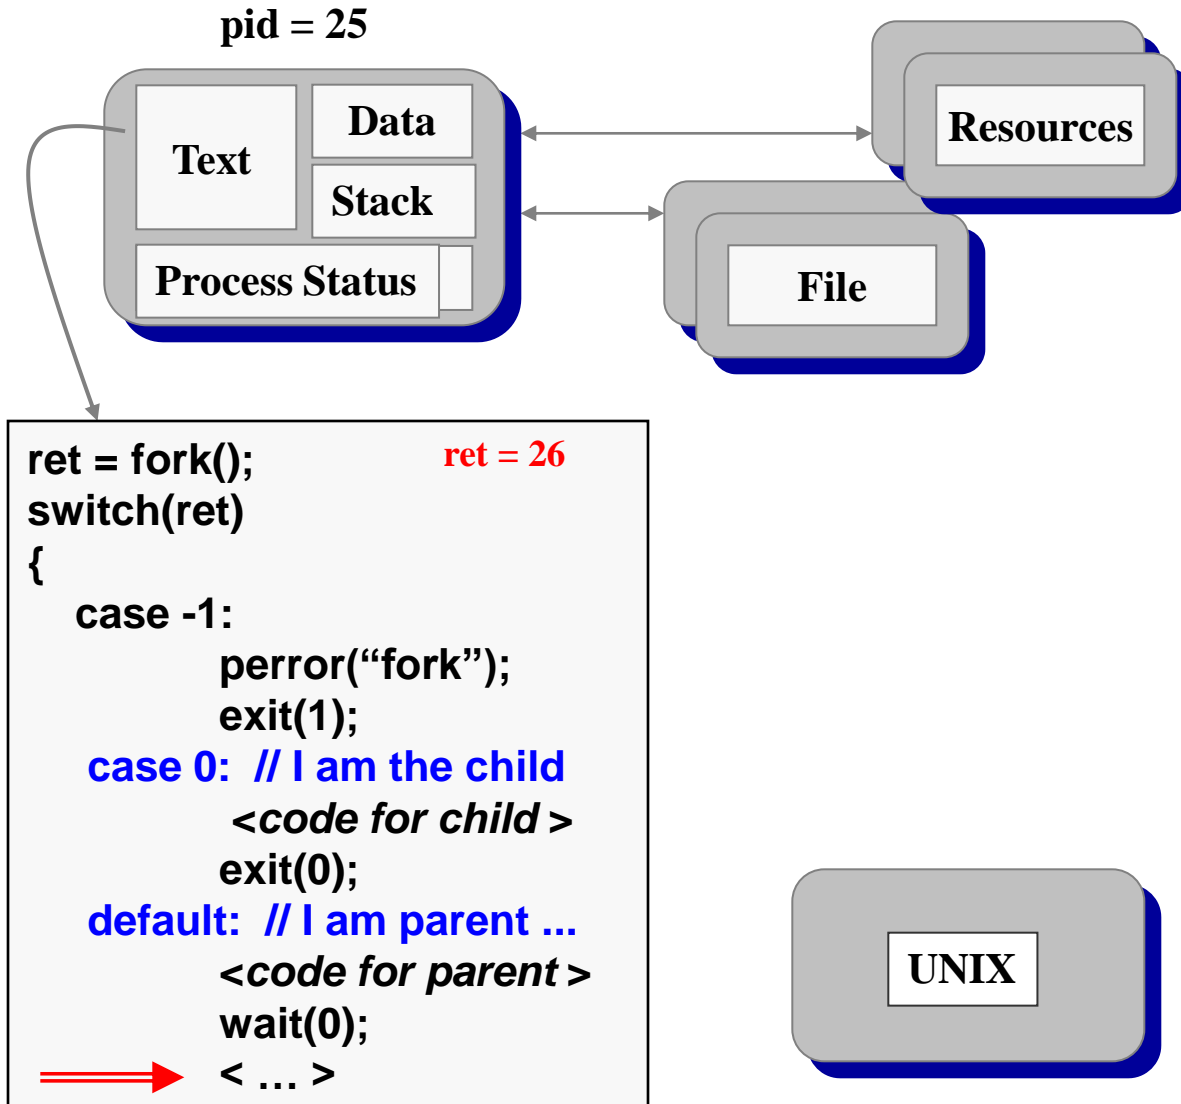
# How fork Works (4)



# How fork Works (5)



# How fork Works (6)



# Execution of a program

---

- If the fork is used to create a new process that must execute a program we use the exec() command
- When executed, the Operating System replaces the current process image (text, data, stack) with a new process image
- The new program must have a main()
- There is no return

In this course we assume that the child is executing a code defined in the parent program

```
ret = fork();
switch(ret)
{
    case -1:
        perror("fork");
        exit(1);
    case 0: // I am the child
        exec*( .... )
    default: // I am parent ...
        <code for parent >
        wait(0);
        < ... >
}
```

# Orderly Termination: `exit()`

---

- To finish execution, a child may call **`exit(status)`**
- This system call:
  - Saves result = argument of `exit`
  - Executes all functions specified with **`atexit(fun)`** and **`on_exit(fun)`**
  - Streams are down loaded with **`fflush()`**
  - Closes all open files, connections
    - (not the ones shared with other processes)
  - Call **`_exit(status)`**

# Orderly Termination: `_exit()`

---

- To finish execution, a child may call `_exit(status)`
- This system call:
  - Saves result = argument of exit
  - Deallocates memory
  - If the process has running childs, they are assigned to init
  - Checks if parent is alive
  - If parent is alive, holds the result value until the parent requests it (with `wait`); in this case, the child process does not really die, but it enters a zombie/defunct state
  - If parent is not alive, the child terminates (dies)

# Waiting for the Child to Finish

---

- Parent may want to wait for children to finish
  - Example: a shell waiting for operations to complete
- Waiting for any some child to terminate: **wait()**
  - Blocks until some child terminates
  - Returns the process ID of the child process
  - Or returns -1 if no children exist (i.e., already exited)
- Waiting for a specific child to terminate: **waitpid()**
  - Blocks till a child with particular process ID terminates

```
#include <sys/types.h>
#include <sys/wait.h>

pid_t wait(int *status);
pid_t waitpid(pid_t pid, int *status, int options);
```

# Processes with Python

---

- Process creation is managed by operating system
- Python must be able to asks the OS to do it
  - `import os`
- Commands are similar
  - `pid = os.fork()`
  - `os.wait()`
  - `os.waitpid()`
  - `os.exit()`
  - `os._exit()`

```
import os

ret = os.fork()
if ret==0:
    # I am the child
    <code for child>
    os._exit(0)
else:
    # I am parent ...
    <code for parent >
    os.wait(0);

< ... >
```





# Roadmap

- Processes: fork(), wait()
- • Threads: resource ownership and execution
- Case study:
  - PThreads
- Symmetric multiprocessing (SMP)





# Processes and Threads

- A process has two characteristics:
  - **Scheduling/execution** - follows an execution path that may be interleaved with other processes
  - **Resource ownership** - includes a virtual address space to hold the process image
- These two characteristics are treated independently by the operating system





# Processes and Threads

- The unit of dispatching is referred to as a ***thread*** or lightweight process
- The unit of resource ownership is referred to as a process or ***task***





# Multithreading

- The ability of an OS to support multiple, concurrent paths of execution within a single process.




Figure 4.1 Threads and Processes [ANDE97]



# Single Thread Approaches

- MS-DOS supports a single user process and a single thread
- Some UNIX support multiple user processes but only support one thread per process




Figure 4.1 Threads and Processes [ANDE97]

# Multithreading

- Often a Java run-time environment is a single process with multiple threads
- Multiple processes **and** threads are found in Windows, Solaris, and many modern versions of UNIX

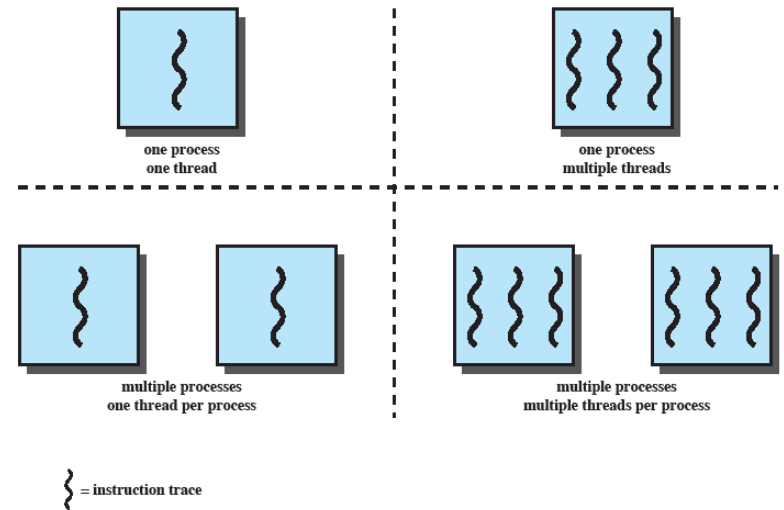
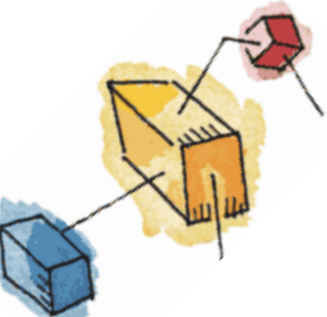


Figure 4.1 Threads and Processes [ANDE97]



# Processes in Multithreaded OS


- A virtual address space which holds the process image
- Protected access to
  - Processors
  - Other processes
  - Files
  - I/O resources





# One or More Threads in Process


- Each thread has
  - An execution state (running, ready, etc.)
  - Saved thread context when not running
  - An execution stack
  - Some per-thread static storage for local variables
  - Access to the memory and resources of its process (all threads of a process share this)





# One view...

- One way to view a thread is as an independent program counter operating *within* a process





# Threads vs. processes

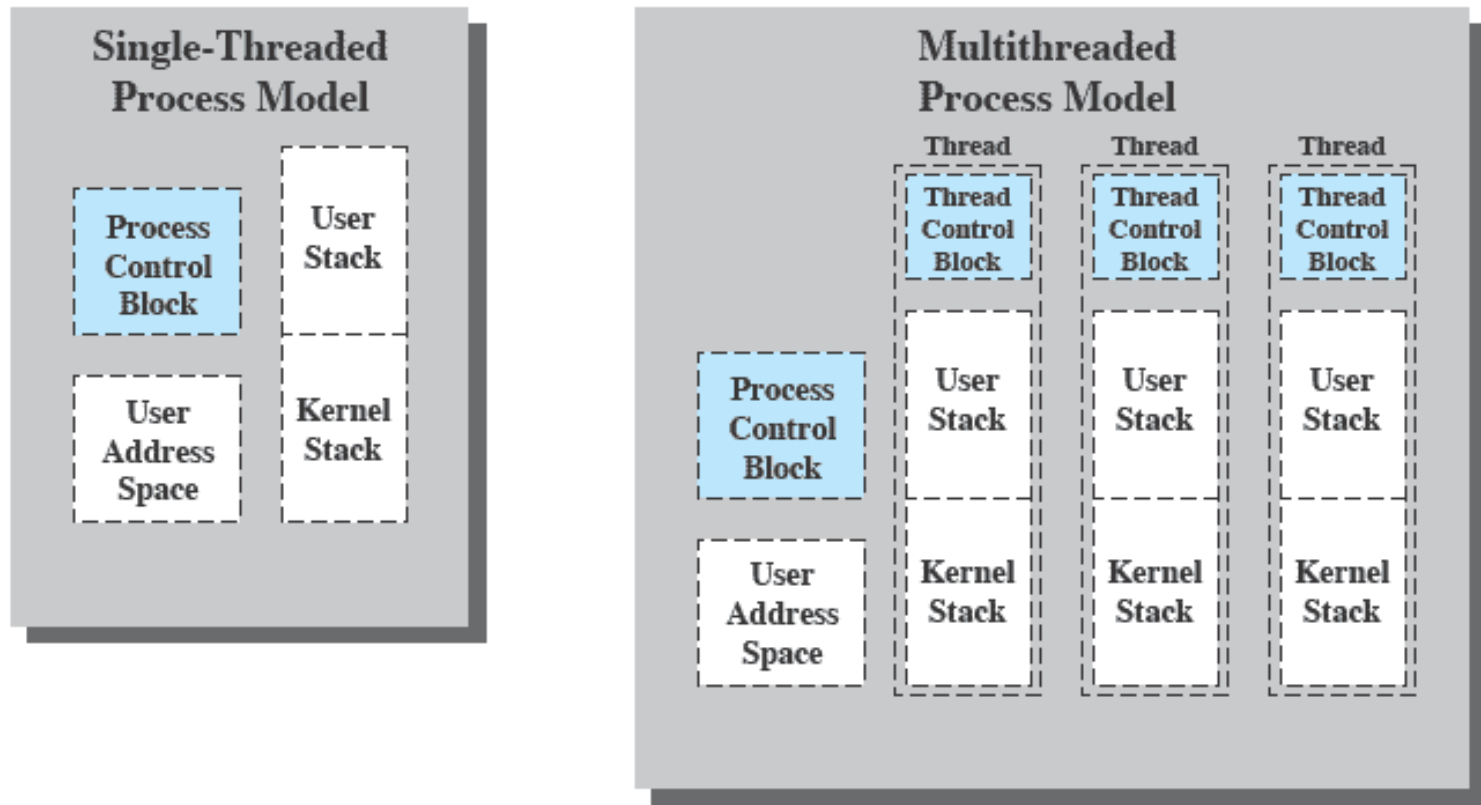
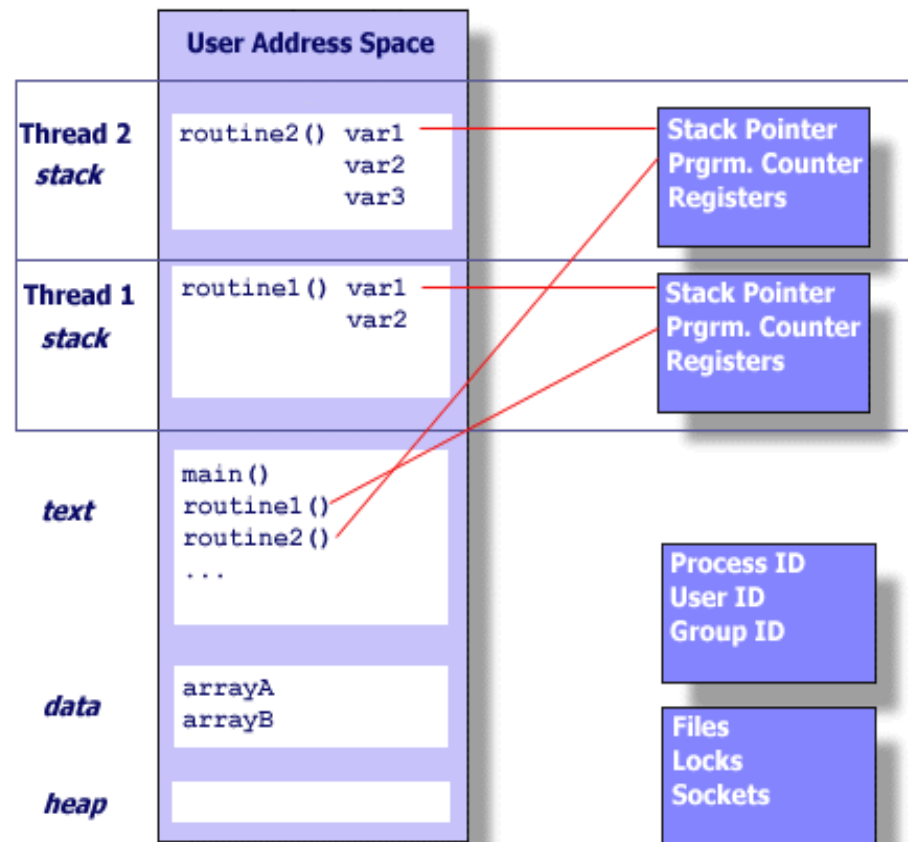
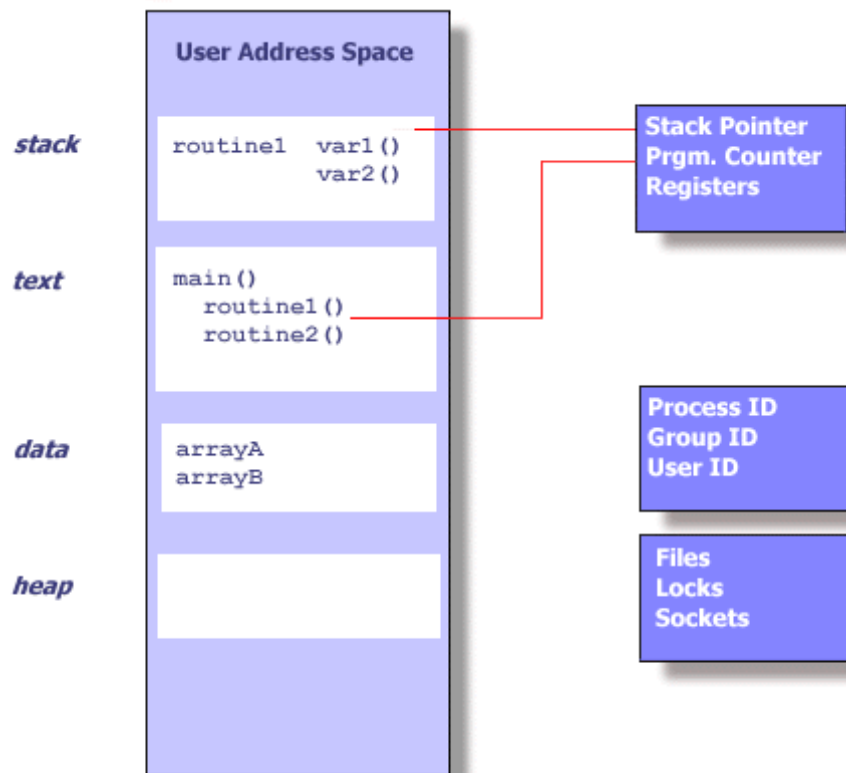



Figure 4.2 Single Threaded and Multithreaded Process Models



# Unix Process vs thread





# Benefits of Threads

- Takes less time to create or terminate a new thread than a process
- Switching between two threads takes less time than switching processes
- Threads can communicate with each other
  - without invoking the kernel





# Thread use in a Single-User System

- Foreground and background work
- Asynchronous processing
- Speed of execution
  - e.g., execution advances while a thread waits for I/O
- Modular program structure





# Threads


- Several actions can affect all of the threads in a process
  - OS must manage these at the process level
- Examples:
  - Suspending a process involves suspending all threads of the process (*same address space!*)
  - Termination of a process terminates all threads within the process



# Activities similar to Processes

- Threads have execution states and may synchronize with one another
  - Similar to processes
- We look at these two aspects of thread functionality in turn
  - States
  - Synchronisation





# Thread Execution States

- States associated to threads:
  - Running, ready, blocked
- To change the thread state
  - Spawn (another thread)
  - Block
    - Issue: can blocking a thread result in blocking some other thread, or even the whole process?
  - Unblock
  - Finish (thread)
    - Deallocate register context and stacks








# Example:

## Remote Procedure Call

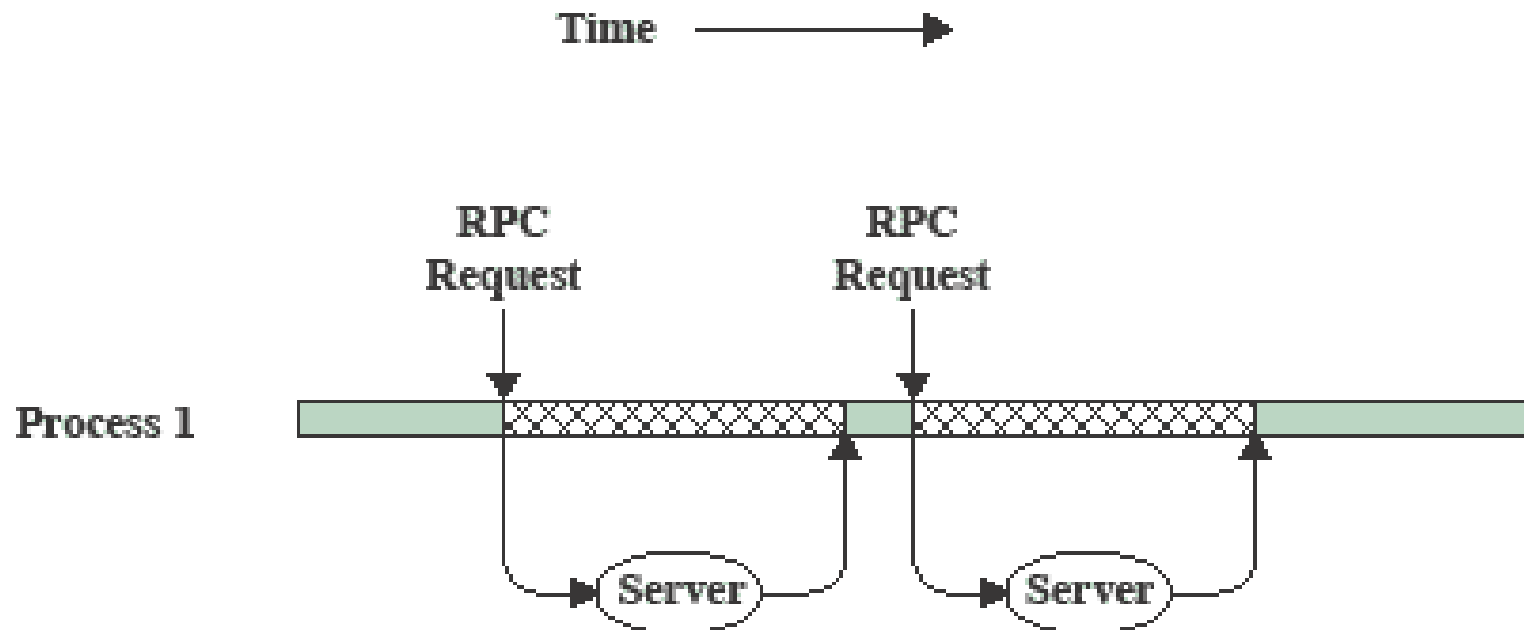
- Consider:
  - A program that performs two remote procedure calls (RPCs)
  - to two different hosts
  - to obtain a combined result.






# RPC

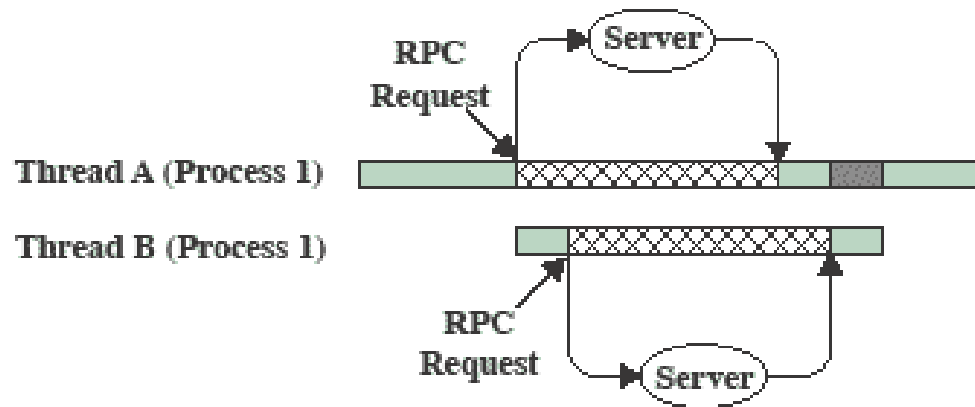
## Using Single Thread






(a) RPC Using Single Thread



# RPC Using One Thread per Server



(b) RPC Using One Thread per Server (on a uniprocessor)

-  Blocked, waiting for response to RPC
-  Blocked, waiting for processor, which is in use by Thread B
-  Running

# Multithreading on a Uniprocessor

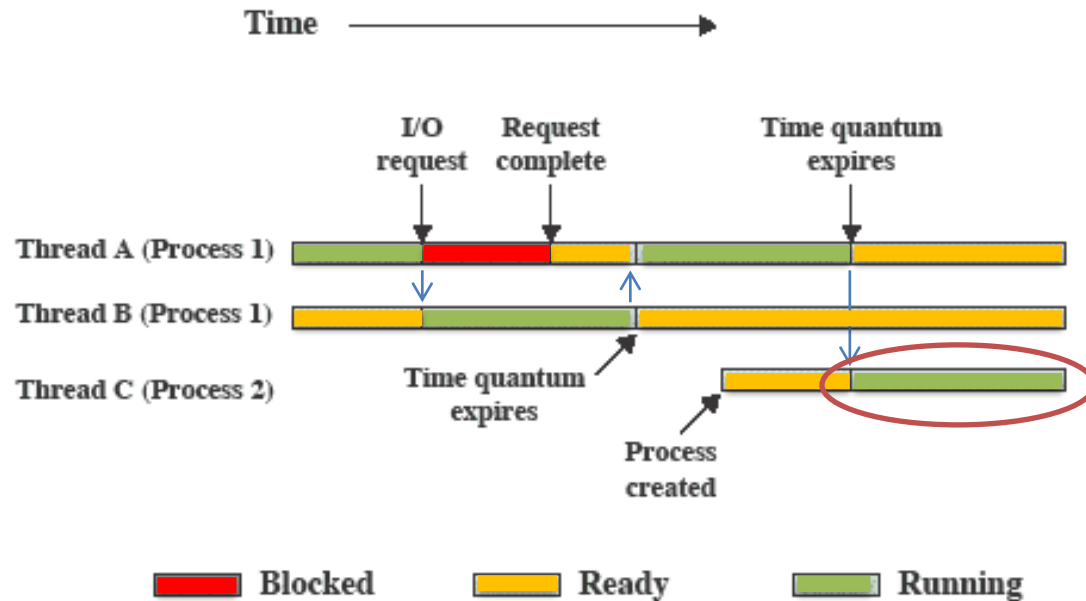


Figure 4.4 Multithreading Example on a Uniprocessor



# Categories of Thread Implementation


- User Level Thread (ULT)
- Kernel level Thread (KLT) also called:
  - kernel-supported threads
  - lightweight processes





# User-Level Threads

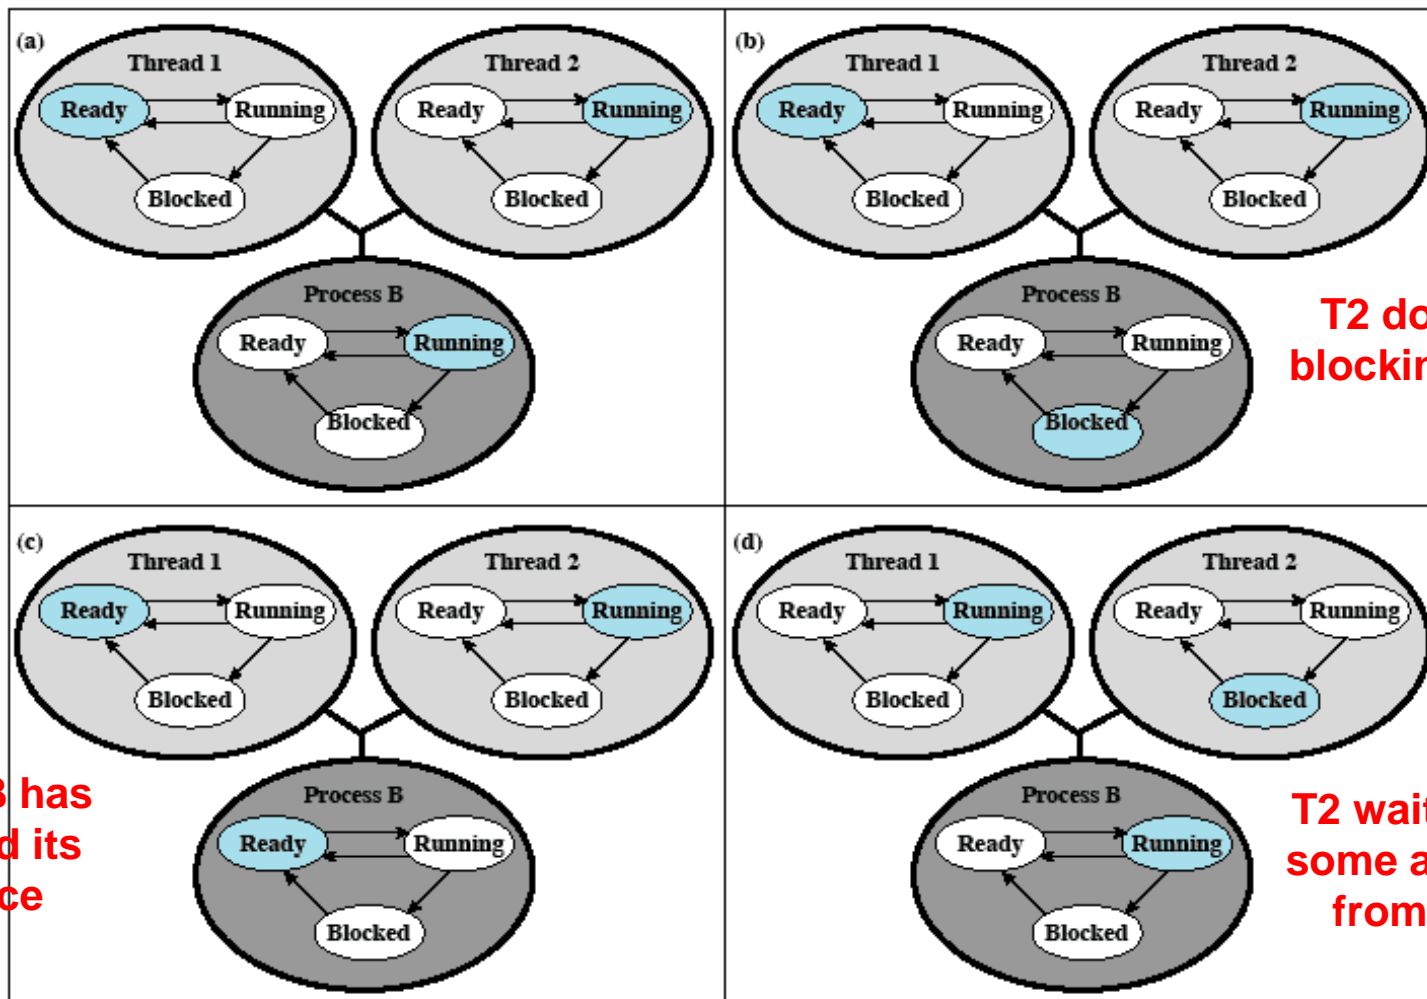
- All thread management is done by the application
- The kernel is not aware of the existence of threads



(a) Pure user-level



# Relationships between ULT Threads and Process States



T2 does a blocking call

Process B has exhausted its time slice

T2 waits for some action from T1


Colored state is current state

Figure 4.7 Examples of the Relationships Between User-Level Thread States and Process States



# Kernel-Level Threads

- Kernel maintains context information for the process and the threads
  - No thread management done by application
- Scheduling is done on a thread basis
- Windows is an example of this approach



(b) Pure kernel-level







# Advantages of ULT

- Application-specific thread scheduling (i.e., independent of kernel)
- Thread switch does not require kernel privilege/switch to kernel mode
- ULTs run on any OS: implementation is done through a thread library at user level





# Disadvantages of ULT

- A blocking systems call executed by a thread blocks *all threads* of the process
- Pure ULTs does not take full advantage of multiprocessors/multicores architectures





# Advantages of KLT

- The kernel can simultaneously schedule multiple threads from the same process on multiple processors
- If one thread in a process is blocked, the kernel can schedule another thread of the same process
- Kernel routines themselves can be multithreaded





# Disadvantage of KLT


- The transfer of control from one thread to another within the same process requires a mode switch to the kernel





# Combined Approaches

- Thread creation done in the user space
- Bulk of scheduling and synchronization of threads done within the application
- $u$  ULTs are mapped onto  $k$  KLTs ( $k=u$  in Solaris)



(c) Combined






# Threads & Processes: Possible Arrangements

Table 4.2 Relationship Between Threads and Processes

Threads:Processes	Description	Example Systems
1:1	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
M:1	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux, OS/2, OS/390, MACH
1:M	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
M:N	Combines attributes of M:1 and 1:M cases.	TRIX





# Roadmap

- Processes: fork (), wait()
- Threads: Resource ownership and execution
- • Case study:
  - PThreads
- Symmetric multiprocessing (SMP).





# POSIX Threads (PThreads)

- For UNIX systems, implementations of threads that adhere to the IEEE POSIX 1003.1c standard are Pthreads.
- Pthreads are C language programming types defined in the `pthread.h` header/include file.







# Why Use Pthreads

- The primary motivation behind Pthreads is improving program performance
- Can be created with much less OS overhead
- Need fewer system resources to run
- Timing comparison (next slide)
  - forking processes vs `pthread_create()`
  - timings reflect 50K process/thread creations (unit: s)





# Threads vs Forks

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
Intel 2.6 GHz Xeon E5-2670 (16 cores/node)	8.1	0.1	2.9	0.9	0.2	0.3
Intel 2.8 GHz Xeon 5660 (12 cores/node)	4.4	0.4	4.3	0.7	0.2	0.5
AMD 2.3 GHz Opteron (16 cores/node)	12.5	1.0	12.5	1.2	0.2	1.3
AMD 2.4 GHz Opteron (8 cores/node)	17.6	2.2	15.7	1.4	0.3	1.3
IBM 4.0 GHz POWER6 (8 cpus/node)	9.5	0.6	8.8	1.6	0.1	0.4
IBM 1.9 GHz POWER5 p5-575 (8 cpus/node)	64.2	30.7	27.6	1.7	0.6	1.1
IBM 1.5 GHz POWER4 (8 cpus/node)	104.5	48.6	47.2	2.1	1.0	1.5
INTEL 2.4 GHz Xeon (2 cpus/node)	54.9	1.5	20.8	1.6	0.7	0.9
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.5	1.1	22.2	2.0	1.2	0.6

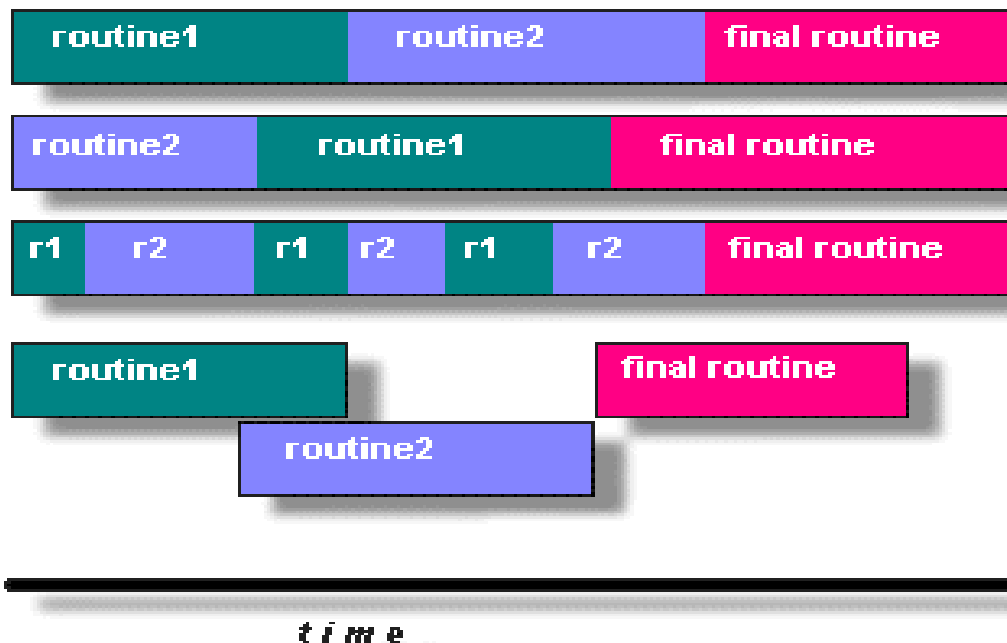
Runtime in seconds to execute 50000 operations





# Designing Threaded Programs as in Parallel Programming

- To take advantage of Pthreads, a program should be organized into discrete, independent tasks that can execute concurrently
  - E.g., if routine1 and routine2 can be interchanged, interleaved and/or overlapped in real time, they are candidates for threading.

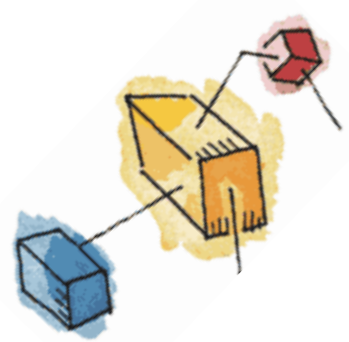




# Models for Threaded Programs

- Manager/worker
  - A *manager* thread assigns work to other threads, the *workers*. Manager handles input and hands out the work to the other tasks
- Pipeline
  - A task is broken into a series of suboperations, each handled in series, but concurrently, by a different thread






# Shared-memory Model

- All threads have access to the *same* global, shared memory
- Threads also have their own private data
- Programmers are responsible for synchronizing access to (i.e., protecting) globally shared data



# Shared-memory Model





# Thread Safety

- A code is thread-safe when multiple threads can execute it simultaneously without unintended interactions
  - (without *clobbering* shared data)
  - (without creating *race conditions*)





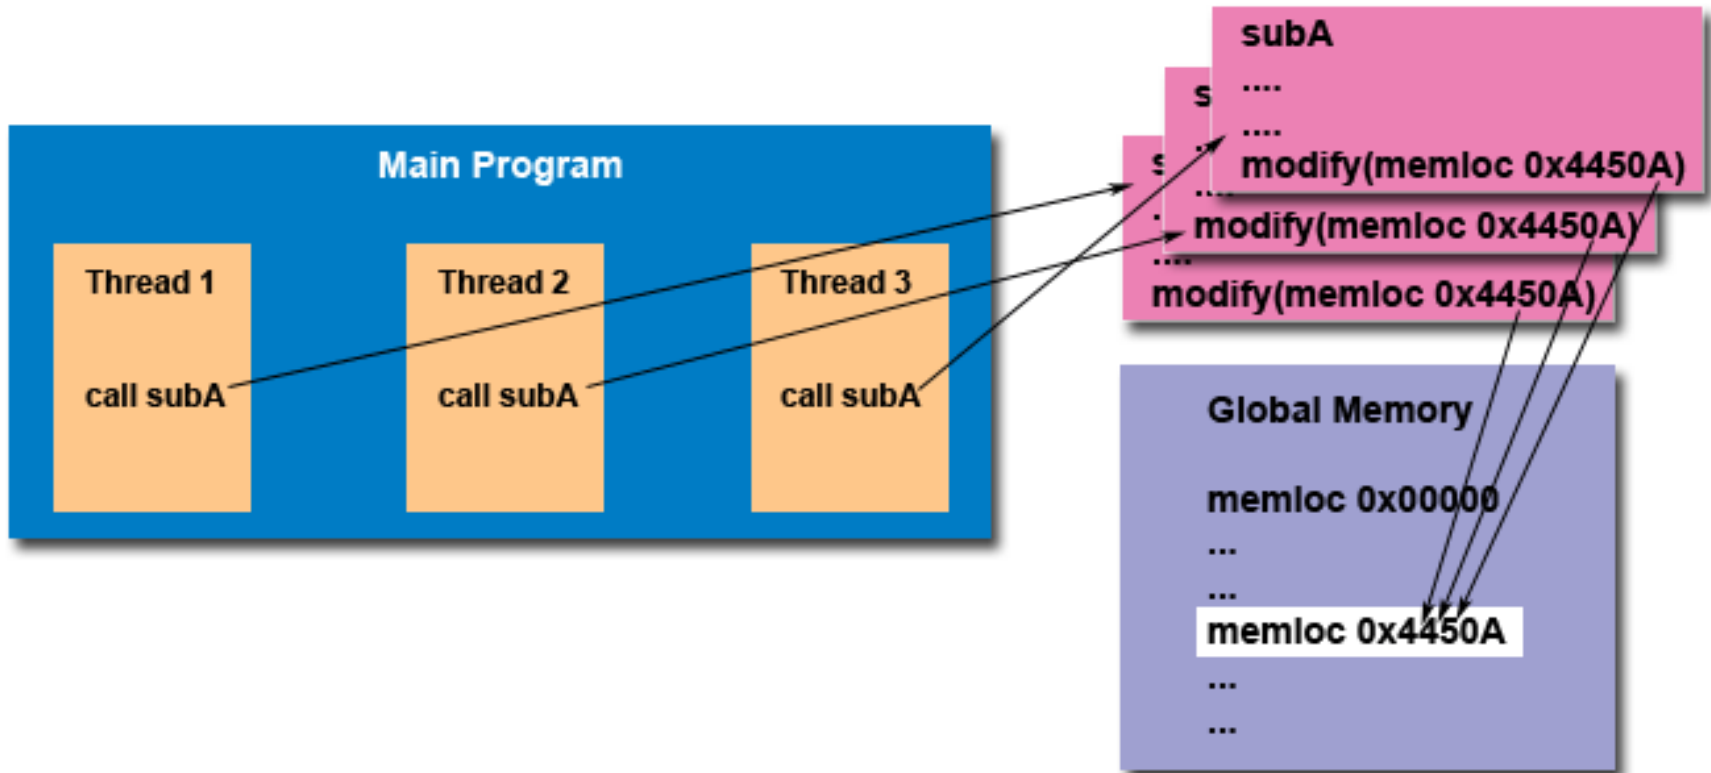
# Thread Safety

- Example: an application creates several threads, each of which makes a call to the same library routine:
  - The library routine accesses/modifies a global structure or location in memory
  - As each thread calls this routine, it is possible that they may try to modify this structure/location at the same time
  - If the routine does not employ some sort of *synchronization mechanism* to prevent data corruption, then it is not thread-safe

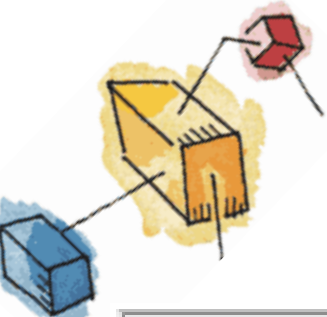




# Thread Safety




# Thread Safety



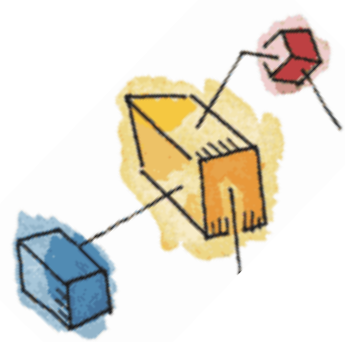
Thread 1	Thread 2	Balance
Read balance: \$1000		\$1000
	Read balance: \$1000	\$1000
	Deposit \$200	\$1000
Deposit \$200		\$1000
Update balance \$1000+\$200		\$1200
	Update balance \$1000+\$200	\$1200

Shared memory



# Pthreads: Creating Threads

- A program's `main()` method comprises a single, default thread.
- `pthread_create()` creates a new thread and makes it executable
  - The maximum number of threads that a process can create is implementation dependent
  - Once created, threads are *peers*, and can create other threads as well





# Pthreads: Terminating Threads

- Several ways to terminate a thread, e.g.:
  - The thread is complete, i.e., the function it started with reaches a `return` statement
  - `pthread_exit()` is called once a thread has completed its work and it is no longer required to exist
  - `pthread_cancel()` from another thread
  - `exit()` is called (*affects the entire program!*)
  - The main terminates without executing `pthread_exit()` [caveat: `pthread_detach()`]





# Pthread: Terminating Threads (cont)

- If the main thread finishes before any other thread does, the other threads will continue executing if `pthread_exit()` was used to terminate the main, or if `pthread_detach()` was used on them
- `pthread_exit()` doesn't free resources (e.g., any file opened inside the thread will stay open), so bear cleanup in mind!




# Pthread Example (1/2)

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS 5

void* printHello(void *arg) {
    int threadID = *(int*) arg;
    printf("Hey! It's me, thread #%d!\n",
          threadID);
    pthread_exit(NULL);
}
```

# Pthread Example (2/2)

```
int main (int argc, char *argv[]) {  
    pthread_t threads[NUM_THREADS];  
    int ret, t;  
    for (t=0; t<NUM_THREADS; t++) {  
        printf("In main: creating thread %d\n", t);  
        int *arg = malloc(sizeof(int));  
        *arg = t;  
        ret = pthread_create(&threads[t], NULL,  
                             printHello, (void*)arg);  
        if (ret != 0) {  
            printf("ERROR: code %d\n", ret); exit(-1);  
        }  
    }  
    pthread_exit(NULL);  
}
```



# One Possible Execution

In main: creating thread 0

In main: creating thread 1

Hey! It's me, thread #0!

In main: creating thread 2

Hey! It's me, thread #2!

Hey! It's me, thread #1!

In main: creating thread 3

In main: creating thread 4

Hey! It's me, thread #3!

Hey! It's me, thread #4!





# Example: Multiple Threads

```
#include <stdio.h>
#include <pthread.h>
#define NUM_THREADS 4

void *hello (void *arg) {
    printf("Hello Thread\n");
}

main() {
    pthread_t tid[NUM_THREADS];
    for (int i = 0; i < NUM_THREADS; i++)
        pthread_create(&tid[i], NULL, hello, NULL);

    for (int i = 0; i < NUM_THREADS; i++)
        pthread_join(tid[i], NULL);
}
```



# Threads in Java

- As a subclass of Thread
  - Or as a class implementing runnable
- Thread has a method `run()` that the subclass must redefine
- We create a new thread with `new()`
- We execute the thread with `start()` (that calls run)
- A thread can wait for another one with a `join()`

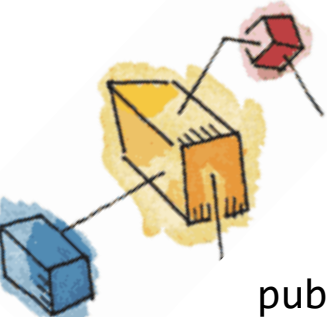




```
public class TwoThreadsTest {  
    public static void main (String[] args) {  
        SimpleThread t1 = new SimpleThread("Jamaica");  
        SimpleThread t2 = new SimpleThread("Fiji");  
        t1.start();  
        t2.start();  
        t1.join();  
        System.out.println("Jamaica thread ended.  
                           I don't care about Fiji thread");  
    }  
}
```



# Example (2)

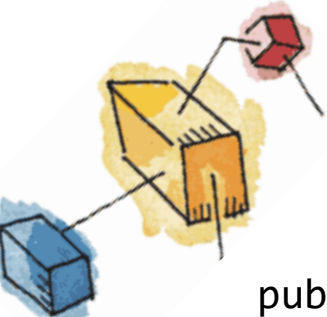


```
public class SimpleThread implements Runnable{
    int i;
    public void run() {
        i=0;
        while (i<10){
            System.out.println(i + "Ciao");
            i++;
        }
    }
}
```

```
public class TwoThreadsTest {
    public static void main (String[] args) {
        SimpleThread s1 = new SimpleThread();
        SimpleThread s2 = new SimpleThread();
        Thread t1 = new Thread(s1);
        Thread t2 = new Thread(s2);
        t1.start();
        t2.start();
    }
}
```



# Example (3)



```
public class SimpleThread implements Runnable{
    int i=0;
    public void run() {
        while (i<10){
            System.out.println(i + "Ciao");
            i++;
        }
    }
}
```

```
public class TwoThreadsTest {
    public static void main (String[] args) {
        SimpleThread s = new SimpleThread();
        Thread t1 = new Thread(s);
        t1.start();
        Thread t2 = new Thread(s);
        t2.start();
    }
}
```

In this case  
the two threads  
are sharing the  
same memory  
and variables



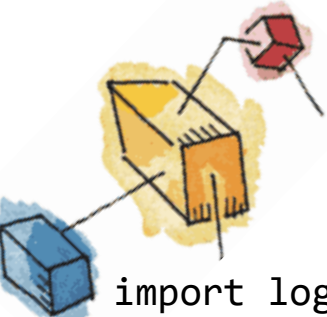


# Thread in Python

- The Python standard library provides threading, which contains most of the needed primitives
- We create a new thread with the method `Thread()`
- We execute the thread with `start()` (that executes the associated function)
- A thread can wait for another one with a `join()`



# Example



```
import logging
import threading
import time
```

```
def thread_function(name):
    logging.info("Thread %s: starting", name)
    time.sleep(2)
    logging.info("Thread %s: finishing", name)
```

```
if __name__ == "__main__":
    format = "%(asctime)s: %(message)s"
    logging.basicConfig(format=format, level=logging.INFO, datefmt="%H:%M:%S")

    logging.info("Main      : before creating thread")
    x = threading.Thread(target=thread_function, args=(1,))
    logging.info("Main      : before running thread")
    x.start()
    logging.info("Main      : wait for the thread to finish")
    x.join()
    logging.info("Main      : all done")
```





# Roadmap

- Processes: fork(), wait()
- Threads: resource ownership and execution
- Case study:
  - Pthreads
- • Symmetric multiprocessing (SMP)







# Traditional View


- Traditionally, the computer has been viewed as a sequential machine
  - A processor executes instructions one at a time in sequence
  - Each instruction is a sequence of operations
- Some popular approaches to parallelism
  - Symmetric MultiProcessors (SMPs)
  - Clusters





# Categories of Computer Systems (Flynn's Taxonomy)


- Single Instruction Single Data (SISD)
  - Single processor executes a single instruction stream to operate on data stored in a single memory






# Categories of Computer Systems


- Single Instruction Multiple Data (SIMD)
  - Each instruction is executed on a different set of data by the different processors






# Categories of Computer Systems


- Multiple Instruction Single Data (MISD) stream
  - A sequence of data is transmitted to a set of processors, each executing a different instruction sequence





# Categories of Computer Systems

- Multiple Instruction Multiple Data (MIMD)
  - A set of processors simultaneously execute different instruction sequences on different data sets



# Parallel Processor Architectures

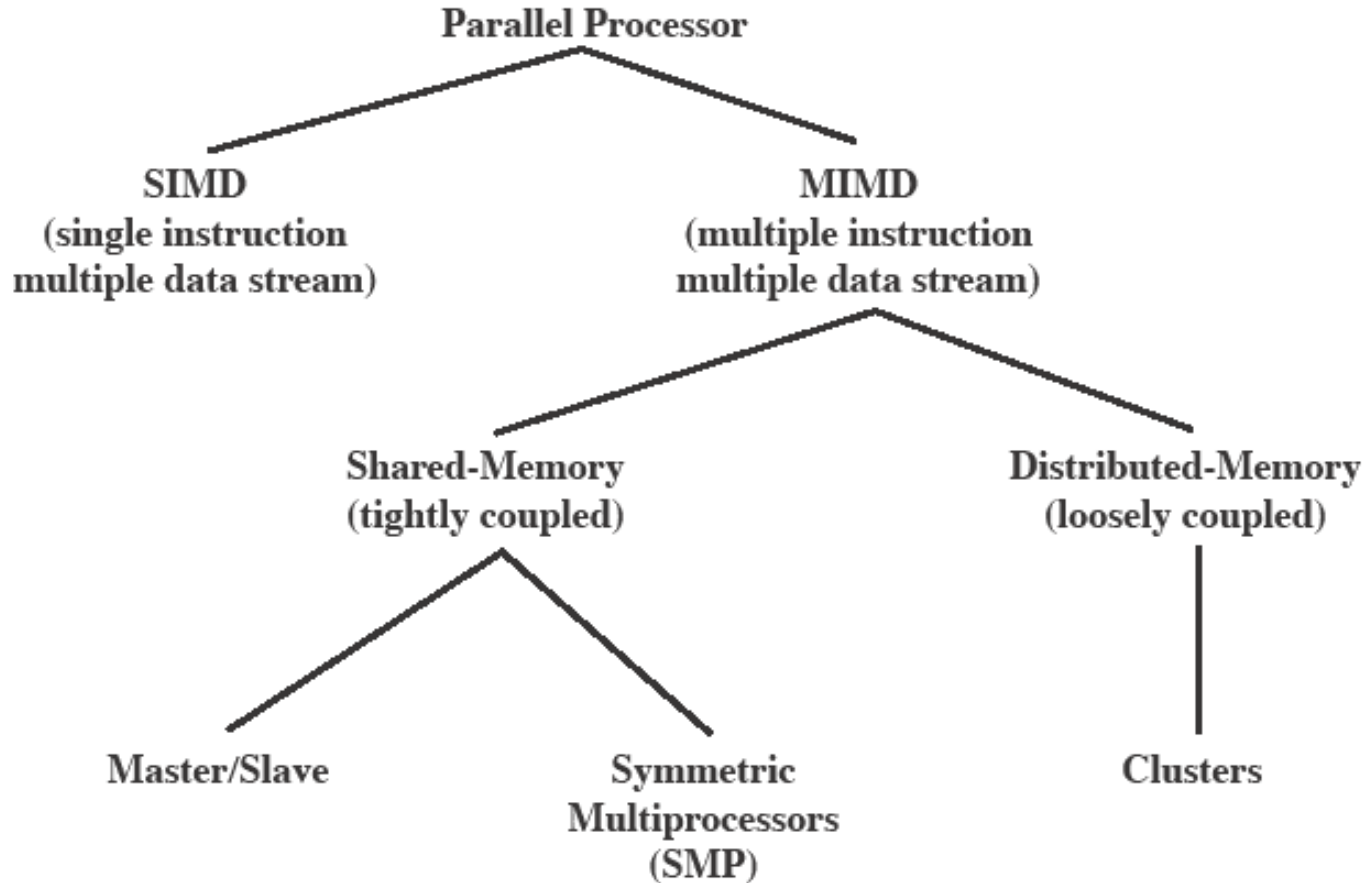


Figure 4.8 Parallel Processor Architectures

# Typical SMP organization

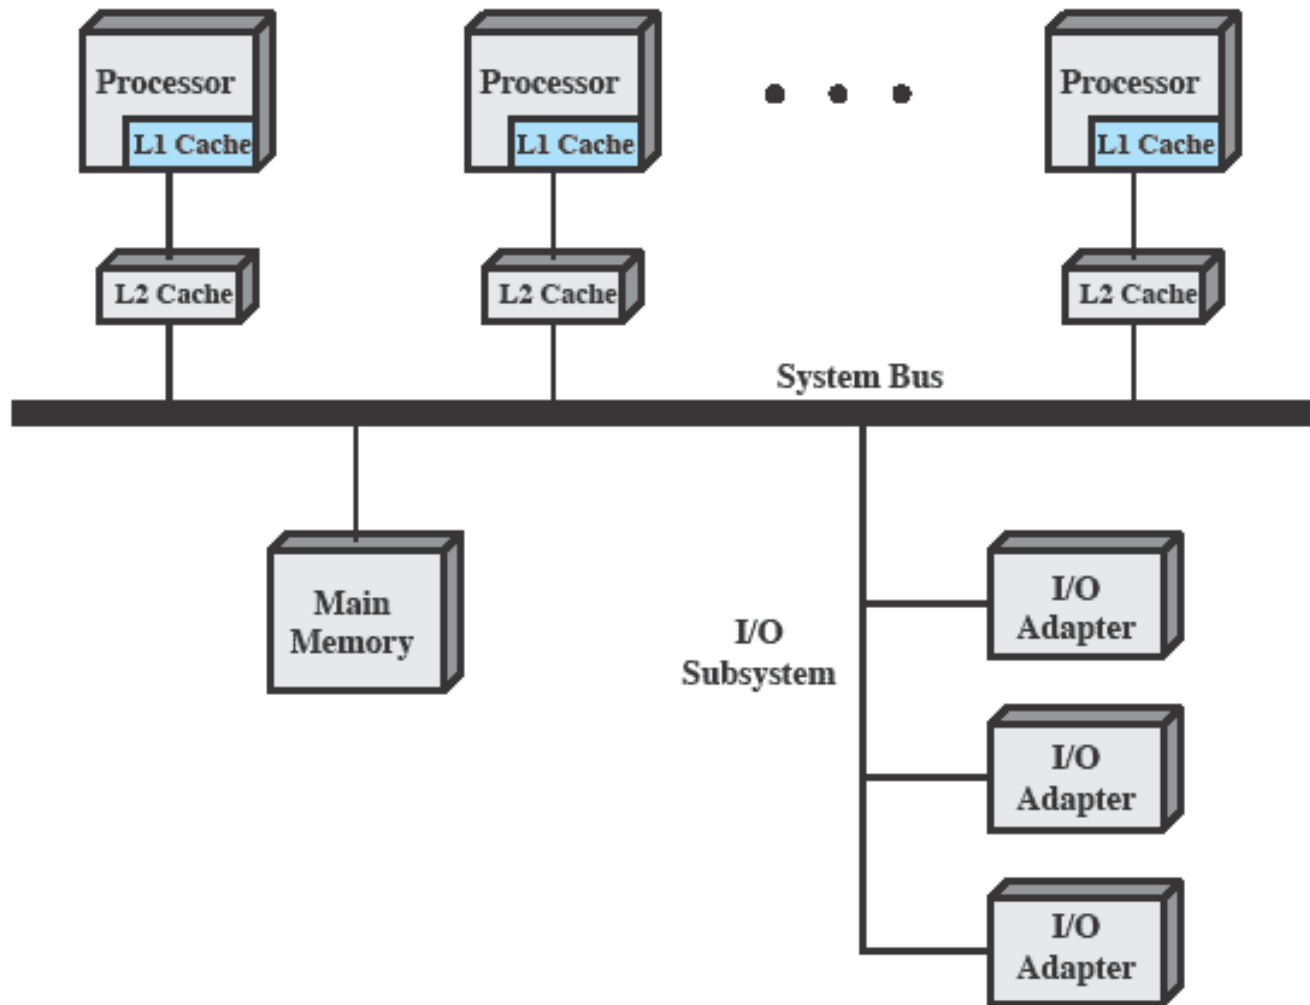
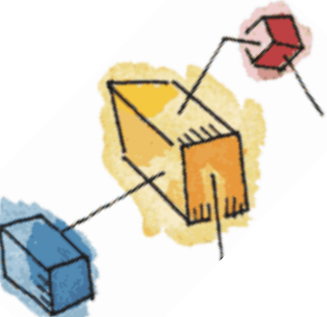


Figure 4.9 Symmetric Multiprocessor Organization



# Multiprocessor OS Design Considerations

- The key design issues include
  - Simultaneous concurrent processes or threads
  - Scheduling
  - Synchronization
  - Memory Management
  - Reliability and Fault Tolerance

